

COMP2005

Histogram Processing

Histograms

- The histogram of a digital image with grey levels in the range $[0, L-1]$ is a discrete function


$$p(r_k) = n_k$$

where:


- r_k is the k th grey level,
 - n_k is the number of pixels in the image with that grey level,
 - $k = 0, 1, \dots, L-1$
- Histograms provide useful global information about the image, ease computation of some image properties, and can be manipulated to improve the image

Histograms

- Dark




- Light




Histograms

- Low Contrast



- High Contrast



Normalised Histograms

- A normalised histogram is a discrete function

$$p(r_k) = n_k / n$$

where:

- n = width x height is the total number of pixels in the image
- The bins in a normalised histogram sum to one
- Each bin gives the probability of the corresponding grey level appearing in the image
- The probabilistic interpretation is valuable in e.g. contrast enhancement and automatic thresholding


Histogram Equalisation

- Goal is to improve the contrast of an image
 - To transform an image in such a way that the transformed image has a nearly uniform distribution of pixel values
 - More general than linear or piecewise contrast stretching
- Histogram transforms
 - Map an input histogram r onto a new histogram s
 - Assume r has been normalized to the interval $[0,1]$, with $r = 0$ representing black and $r = 1$ representing white

$$s = T(r) \quad 0 \leq r \leq 1$$

Histogram Equalisation


- The transformation function satisfies the following conditions
 - $T(r)$ is single-valued and strictly monotonically increasing in the interval $0 \leq r \leq 1$
 - $0 \leq T(r) \leq 1$ for $0 \leq r \leq 1$
- This means it is possible to invert the process
- It also gives us the relationship that allows the derivation of histogram equalisation



Histogram Equalisation

$$P(y).dy = P(x).dx$$

$$\text{So } P(y) = P(x) \cdot \frac{dx}{dy}$$



Histogram Equalisation

- In Gonzalez and Woods' notation....
- Let $p_r(r)$ and $p_s(s)$ denote the probability density function of random variables r and s , respectively
- If $p_r(r)$ and $T(r)$ are known, then the probability density function $p_s(s)$ of the transformed variable s can be obtained

$$p_s(s) = p_r(r) \left| \frac{dr}{ds} \right|$$

- $r = T^{-1}(s)$

Histogram Equalisation

- If we choose as the transformation function the cumulative distribution function or CDF:

$$T(r) = \int_0^r p_r(w)dw$$

$$\frac{ds}{dr} = \frac{dT(r)}{dr} = \frac{d}{dr} \int_0^r p_r(w)dw = p_r(r)$$

$$p_s(s) = p_r(r) \left| \frac{dr}{ds} \right| = p_r(r) \left| \frac{1}{p_r(r)} \right| = 1 \quad 0 \leq s \leq 1$$

$T(r)$ depends on $p_r(r)$, but the resulting $p_s(s)$ is always uniform.

Histogram Equalisation

- We have a discrete histogram, not a PDF of a continuous random variable
- The probability of occurrence of gray level r_k in an image is

$$p_r(r_k) = \frac{n_k}{n} \quad k = 0, 1, 2, \dots, L - 1$$

- The transformation function is

$$s_k = T(r_k) = \sum_{j=0}^k p_r(r_j) = \sum_{j=0}^k \frac{n_j}{n} \quad k = 0, 1, 2, \dots, L - 1$$

- An output image is obtained by mapping each pixel with level r_k in the input image into a corresponding pixel with level s_k .

In Practice


- To perform histogram equalisation
 - Compute the CDF of the input image.
 - For each pixel in the input image, the corresponding output pixel intensity is calculated by using the CDF as a look-up function.
 - CDF values will be in the range 0 - 1, scale the equalised image to fit the range supported by the output image format.
- The histogram of the output image will be approximately uniform

In Practice


- Consider a 64 x 64 pixel, 3 bit (8 grey level) image
(Example from Gonzalez & Woods)

r_k	n_k	$P_r(r_k)$	$T(r_k)$	s_k	Round
0	790	0.19	0.19	1.33	1
1	1023	0.25	0.44	3.08	3
2	850	0.21	0.65	4.55	5
3	656	0.16	0.81	5.67	6
4	329	0.08	0.89	6.23	6
5	245	0.06	0.95	6.65	7
6	122	0.03	0.98	6.86	7
7	81	0.02	1.00	7.00	7


In Practice



Input histogram




$T(r_k)$, scaled back to 0 - 7





Equalised histogram

Results

■ Light




■ Dark






Results

- Low contrast




- High Contrast



Results

- Each of the four transformations above used a different transform, tuned to the input histogram



Strengths & Weaknesses

- Histogram equalisation works well when the input images
 - aren't too noisy
 - don't have large bright or dark areas





Strengths & Weaknesses

- Here the bright sky has dominated the process, equalisation has introduced an artificial boundary between sunlight and sky but not enhanced the three people



Strengths & Weaknesses

- Here the bright area of interest is enhanced, but the noise in the upper dark region is also more obvious




Medical Applications (MRI)


Input Image



Original Histogram before equalization




Original Cumulative Distribution Function



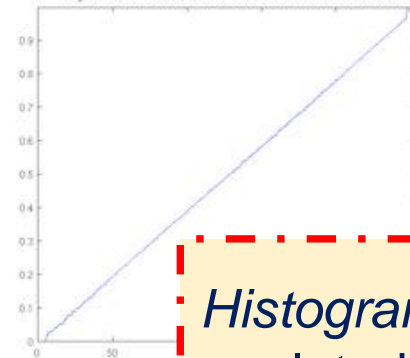
Output Image



Equalized histogram



Equalized Cumulative Distribution Function




Histogram Specification is a related method which transforms an image's histogram so that it matches a target histogram


More Generally...

- *Histogram Specification* is a related method which transforms an image's histogram so that it matches a target histogram


Input
histogram




Target
histogram



Transform



Result



Conclusion

- Histograms, particularly normalised histograms, provide useful summary information
- Images can be manipulated by manipulating their histograms
- Histogram equalisation is a powerful and widely used image enhancement operation
- Global equalisation has drawbacks which can be addressed using local processing
- The method generalises to Histogram Specification

PART 2 : APPLICATIONS

Image Matching with Colour Histograms

Storing & Retrieving Images

- Given a large image database, find all the images containing e.g. horses
- We will focus on individual images, but many of the problems & methods discussed extend to video databases



Text-based Approaches

- Annotation: Relevant words/phrases are added to each image
- Retrieval is via text search
- But annotation is
 - subjective
 - laborious
 - unnecessary????



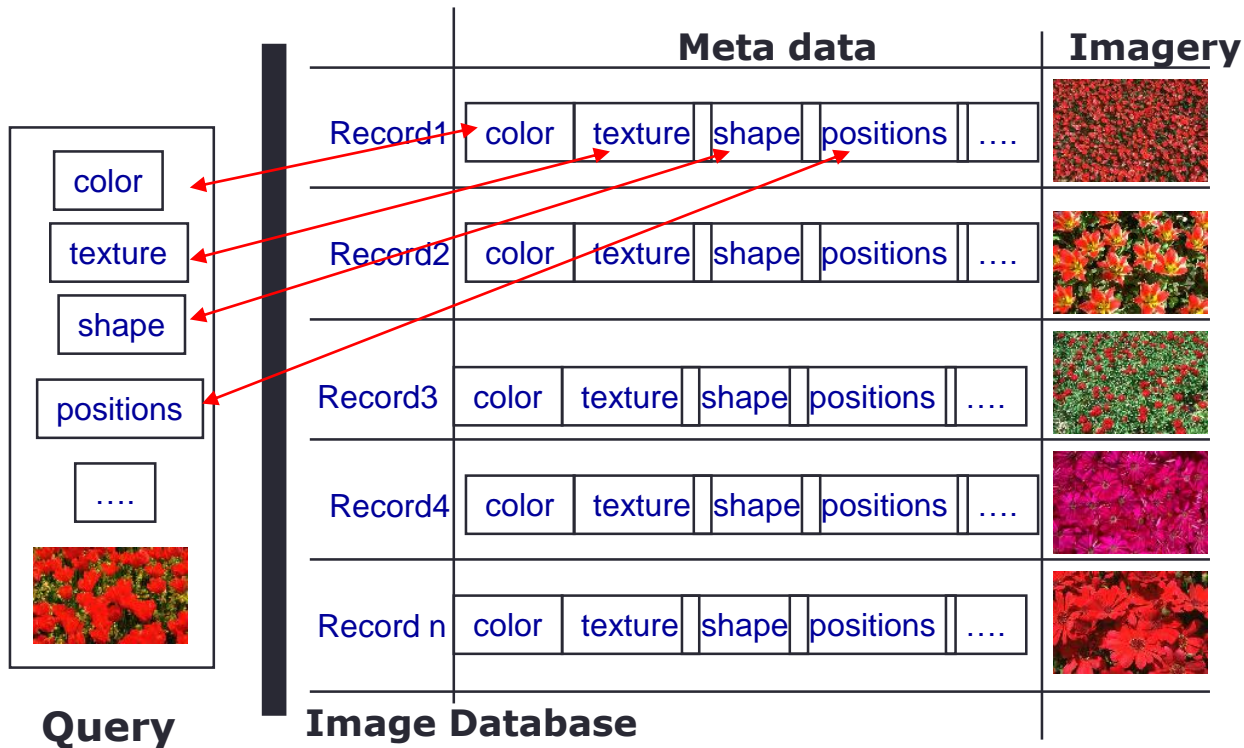
Mother, Child, Vegetable,
Yellow, Green, Purple

Content-based Retrieval

- Indexes the image database on visual features
 - colour
 - shape
 - Texture
- Queries are expressed in those terms or via visual examples
- Simple approaches compute metric distances between the query image and each image in the database
- Advanced approaches use AI techniques, machine learning, etc., and may be interactive
- Some simple measures and datastructures can be very useful




Content-based Retrieval

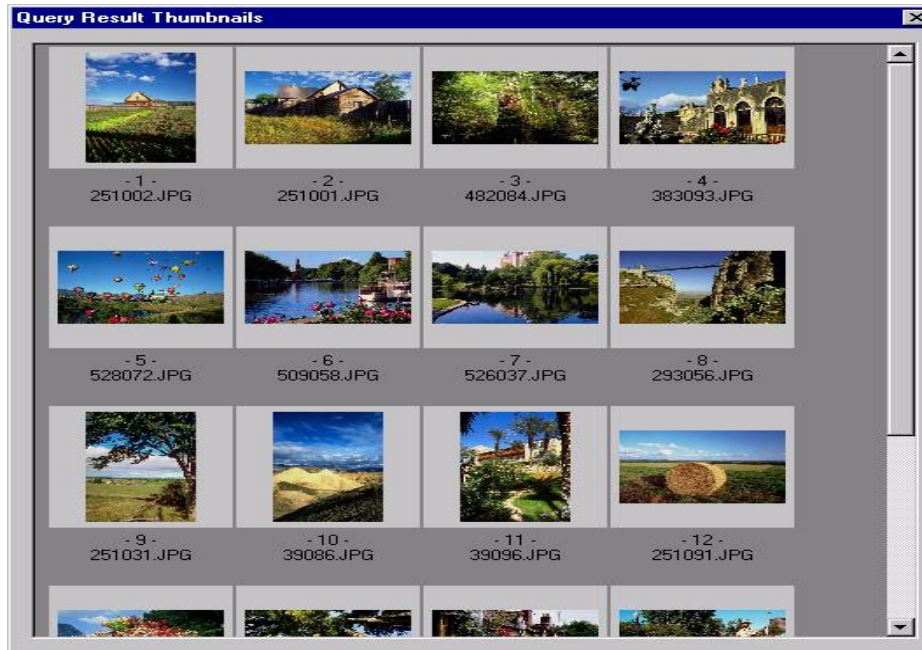


Colour Histograms

- Choose a colour space
 - RGB, HSV,,...
- Divide the axes to create a reasonable number of divisions
 - Trade-off detail against memory/computational cost
- Build a histogram
- Normalise if images are different sizes or colour resolution



Why Colour Histograms?



- Images with similar colour distributions look similar:
colour distribution = colour histogram

Why Colour Histograms?



- Colour correlates well with class identity



- Human vision works hard to preserve colour constancy: presumably because colour is useful

- Histograms
 - Are invariant to translation and rotation
 - Change slowly as viewing direction changes
 - Change slowly with object size
 - Change slowly with occlusion
- Colour histograms summarise target objects quite well, and should match a good range of images

Colour Histograms



- Points in a high dimensional space or compact representations of images?

Common Distance Metrics

- Euclidean or straight-line distance or L2-norm, D^2

$$D^2(H^1, H^2) = \sqrt{\underbrace{\sum_i (H_i^1 - H_i^2)^2}_{\text{Root-mean square error}}} = \|H^1 - H^2\|_2$$

- City-block distance or L1-norm, D^1


$$D^1(H^1, H^2) = \sum_i \underbrace{|H_i^1 - H_i^2|}_{\text{sum of absolute differences}} = \|H^1 - H^2\|_1$$

- *How far apart are two points?*


Some Problems

- Colour quantisation

- Noise and/or different camera responses can give similar images very different histograms



(0,40,0,0,0,0,0,110,0)



(0,0,40,0,0,0,0,0,110)

- Histogram resolution

- May need many bins (4096) to accurately store colour distributions


- Expensive

- The illumination may be coloured

- Same object may generate a different histogram under different lighting

Some Problems

- Colour histograms ignore spatial information
- More advanced methods take spatial relationships into account



- But the comparatively simple image processing methods and representations methods covered so far can do useful things

Histogram Intersection

- Measures how much of the query may be present in the target image (and vice-versa)
- A bin in the target histogram can have a larger value than the corresponding query bin (and vice-versa)

$$HI(H^1, H^2) = \sum_{i=1}^n \min(H_i^1, H_i^2)$$

$H^1 = (10, 0, 0, 0, 100, 10, 30, 0, 0)$
 $H^2 = (0, 40, 0, 0, 0, 6, 0, 110, 0)$
 $HI(H^1, H^2) = 0 + 0 + 0 + 0 + 0 + 6 + 0 + 0 = 6$

- How much do two representations overlap?*

Histogram Intersection

- In the first histogram intersection paper (Ballard and Swain 1991):
 - A database of 66 colour histograms
 - 32 query images
- Recognition rate was almost 100%



?





Histogram Intersection

- Ballard and Swain used opponent colour axes

- $RG = R - G$
- $BY = 2 * B - R - G$
- $WB = R + G + B$ (intensity)

- & matched images under a range of conditions

- Normal condition
- Varying in view
- Varying in image resolutions
- Occlusion (of bottom 1/3 and/or side 1/3 of image)
- Varying in bin resolutions
- Varying in light intensity



Histogram Intersection

Condition	Placement			
	1st	2nd	3rd	>3rd
Full size	29	3	0	0
128 x 90	29	3	0	0
64 x 45	27	5	0	0
32 x 22	14	7	1	0
16 x 11	15	6	4	0
8 x 5	4	4	3	21
Bottom occluded	27	4	1	0
Bottom & side occluded	22	6	5	0

Good until
images are
very small




Not bad when
object only
partly visible



Using Histograms: Object Location


Object Location

- Matching whole images isn't always appropriate.....
 - e.g. if the target object is only expected to fill part of the image or you want to know where it is



Region/Object-based Queries

- We know the rabbit is mostly grey
- Divide the image into windows and see how grey each window is
 - Highlight pixels in the image that are similar to those in the query
 - Look for regions with lots of these pixels



✓ Similar colours

X Dissimilar colours

The Histogram Ratio



- BUT: the image is usually much bigger than the query region

$$R_j = \min(\frac{M_j}{I_j}, 1)$$


- Other objects in the image may be the same colour as parts of the query object
- So some colours are not reliable cues: if you're looking for a zebra at night look for white pixels
- Compute the ratio of corresponding Model and Image histogram bins
- If the image has many more pixels of a given colour R_j is small and that colour is not useful
- If the model has more, R_j is 1 and that colour is useful

Backprojection

- The greater the value of R_j , the more valuable the colour(s) represented by bin j
 - Consider each image pixel
 - If that pixel maps to histogram bin k , replace the pixel value with a grey value $= R_k$
- This is still image processing: we've processed an image to create an image
 - The pixel values are related to the likelihood of their showing the target



Backprojection



- Regions with high average brightness are likely to contain the rabbit – its more complex, but still a histogram matching approach (see Ballard and Swain)

We've seen Backprojection before

- Skin shows a **very** clear peak and narrow range in U,V

