

COMP2005

Compression

Image Compression

- Types of redundancy:
 - Coding redundancy
 - Spatial redundancy
 - Psychovisual redundancy
- Structure of compression systems
- Components and complete schemes:
Huffman coding, GIF, JPEG



Why?

- Compression affects image quality
- Need to understand/choose methods/parameters
- New methods are based on core concepts

Coding Redundancy

- The grey level histogram of an image gives the probability (frequency) of occurrence of grey level r_k

$$p(r_k) = \frac{n_k}{n} \quad k = 0, 1, \dots, L-1$$

- If the number of bits used to represent each value of r_k is $l(r_k)$, the **average number of bits required to represent a pixel** is

$$L_{avg} = \sum_{k=0}^{L-1} l(r_k) p(r_k)$$

- To code an $M \times N$ image requires MNL_{avg} bits

Coding Redundancy

- If an m-bit natural binary code is used to represent grey level then
 - all pixels take the same amount of space,
 - $P(r_k)$ values sum to 1, so

$$L_{avg} = \sum_{k=0}^{L-1} l(r_k) p(r_k) = \sum_{k=0}^{L-1} m p(r_k) = m$$

- and an image occupies MNm bits
- But some pixel values are more common than others.....

Coding Redundancy : Variable Length Encoding

- Assigning fewer bits to the more probable grey levels than to less probable ones can achieve data compression, e.g:

r_k	$p_r(r_k)$	Code 1	$l_1(r_k)$	Code 2	$l_2(r_k)$
$r_{87} = 87$	0.25	01010111	8	01	2
$r_{128} = 128$	0.47	10000000	8	1	1
$r_{186} = 186$	0.25	11000100	8	000	3
$r_{255} = 255$	0.03	11111111	8	001	3
r_k for $k \neq 87, 128, 186, 255$	0	—	8	—	0

- Build a codebook, replace 'true' pixel values with code
- Lossless: the process can be reversed by inverting the codebook

Huffman Coding : Entropy

- The idea: associate information with probability
- A random event E with probability $P(E)$ contains:

$$I(E) = \log\left(\frac{1}{P(E)}\right) = -\log(P(E))$$

units of information

- **Suppose that grey level values are generated by a random variable, then r_k contains:**


$$I(r_k) = -\log(P(r_k))$$

Note: $I(E)=0$ when $P(E)=1$

units of information



Entropy

- Entropy is the average information content of an image, a measure of histogram dispersion



$$H = -\sum_{k=0}^{L-1} p(r_k) \log_2 p(r_k)$$

entropy=7.4635



entropy=0

- Can't compress to less than H bits/pixel without losing information

Huffman Coding

- Compute probabilities of each symbol by histogramming the source
- Process probabilities to pre-compute codebook:
code(i)
 - codebook is static (fixed)
- Encode source symbol-by-symbol:
symbol(i) \rightarrow code(i)
- Transmit coded signal and codebook
- The need to pre-process (histogram) the source before encoding begins is a disadvantage

Huffman Coding

- Builds a binary tree in which symbols to be coded are nodes
- The algorithm:
 - Create a list of nodes, one per symbol, sorted in order of symbol frequency (or probability)
 - REPEAT (until only one node left)
 - Pick the two nodes with the lowest frequencies/probabilities and create a parent of them
 - **Randomly** assign the codes 0,1 to the two new branches of the tree and delete the children from the list
 - Assign the sum of the children's probabilities to their parent and insert it in the list
- Path from root to node gives code for corresponding symbol

Huffman Coding


Original source		Source reduction			
Symbol	Probability	1	2	3	4
a_2	0.4	0.4	0.4	0.4	0.6 0.4
a_6	0.3	0.3	0.3	0.3	
a_1	0.1	0.1	0.2	0.3	
a_4	0.1	0.1	0.1		
a_3	0.06	0.1			
a_5	0.04				

- Create parent node of a_3 and a_5

- Create parent of new node and a_4

- New parent is higher up list (0.2) than a_1 (0.1)

Huffman Coding



Huffman Coding

- The algorithm systematically places nodes representing high probability symbols further up the tree: their paths (and so codes) are shorter
- No code is a prefix to any other – don't need to mark boundaries between codes
 - e.g. 01101010 must be a_1a_3
- In this example
 - Average length of the code is 2.2.bits/symbol
 - The entropy of the source is 2.14 bits/symbol
- In general
 - Break image into small (e.g. 8 x 8) blocks
 - Each block is a symbol to be encoded

A Huffman Code Example

From a past exam paper:

An image has the following normalized histogram. Derive a Huffman code for each pixel value, showing how you obtained your code

10 mins

Pixel Value	Normalised Frequency
0	0.1
1	0.1
2	0.15
3	0.35
4	0.2
5	0
6	0.05
7	0.05

A Huffman Code Example

Pixel Value	Normalised Frequency
0	0.1
1	0.1
2	0.15
3	0.35
4	0.2
5	0
6	0.05
7	0.05

Sort



Pixel Value	Normalised Frequency
3	0.35
4	0.2
2	0.15
1	0.1
0	0.1
7	0.05
6	0.05
5	0

A Huffman Code Example: Table View

Pixel Value	Normalised Frequency							
3	0.35	0.35	0.35	0.35	0.35	0.4	0.6	1.0
4	0.2	0.2	0.2	0.2	0.25	0.35	0.4	
2	0.15	0.15	0.15	0.2	0.2	0.25		
1	0.1	0.1	0.1	0.15	0.2			
0	0.1	0.1	0.1	0.1				
7	0.05	0.05	0.1					
6	0.05	0.05						
5	0							

The diagram illustrates the merging process of Huffman coding. Red arrows show the combination of frequencies from left to right. For example, 0.05 and 0.05 merge to 0.1, then 0.1 and 0.1 merge to 0.2, and so on, until the total frequency of 1.0 is reached.

A Huffman Code Example: The Code

		Code length	Normalised Freq,
3: 00	0: 110	3	0.1
4: 10	1: 011	3	0.1
2: 010	2: 010	3	0.15
1: 011	3: 00	2	0.35
0: 110	4: 10	2	0.2
7: 1110	5: 11111	5	0
6: 11110	6: 11110	5	0.05
5: 11111	7: 1110	4	0.05

Mean bits/pixel $L_{avg} = 0.3 + 0.3 + 0.45 + 0.70 + 0.4 + 0 + 0.25 + 0.2 = 2.6$

$$L_{avg} = \sum_{k=0}^{L-1} l(r_k) p(r_k)$$

Without compression $L_{avg} = 3$

Compression ratio = $2.6/3 = 0.86$

$l(r_k)$: Number of bits used to represent each value of r_k is

L_{avg} : **average number of bits required to represent a pixel** is

Spatial Redundancy

- Sometimes called *Interpixel Redundancy*
- Neighbouring pixels often have similar values
- Compression based on spatial redundancy involves some element of **pixel grouping, or transformation**
- Simplest is ***Run-Length Encoding***
 - maps the pixels along each scan line into a sequence of pairs $(g_1, r_1), (g_2, r_2), \dots,$
 - where g_i is the i th grey level, r_i is the run length of i th run

A Binary Example

Row 1:	(0, 16)	
Row 2:	(0, 16)	
Row 3:	(0, 7) (1, 2) (0, 7)	encode
Row 4:	(0, 4), (1, 8) (0, 4)	←
Row 5:	(0, 3) (1, 2) (0, 6) (1, 3) (0, 2)	
Row 6:	(0,2) (1, 2) (0,8) (1, 2) (0, 2)	
Row 7:	(0, 2) (1,1) (0, 10) (1,1) (0, 2)	
Row 8:	(1, 3) (0, 10) (1,3)	
Row 9:	(1, 3) (0, 10) (1, 3)	
Row 10:	(0,2) (1, 1) (0,10) (1, 1) (0, 2)	
Row 11:	(0, 2) (1, 2) (0, 8) (1, 2) (0, 2)	
Row 12:	(0, 3) (1, 2) (0, 6) (1, 3) (0, 2)	
Row 13:	(0, 4) (1,8) (0, 4)	
Row 14:	(0, 7) (1, 2) (0, 7)	
Row 15:	(0, 16)	decode
Row 16:	(0, 16)	

Psychovisual Redundancy – e.g. GIF compression

- Some grey level and colour differences are imperceptible; goal is to compress without noticeable change to the image

256 gray levels



16 gray levels



16 gray levels



A simple method:
add a small
random number to
each pixel before
quantization

GIF : Vector Quantisation

- Palettised images (gif)
 - Map vector values (R,G,B) onto scalar values
 - Multiple vectors map to each scalar




Vector
quantisation




True colour R,G,B
8 bits each
1677216 possible colours

gif
8 bits per pixel
256 possible colours


Paletised Images




Building a Palette



Building a Palette



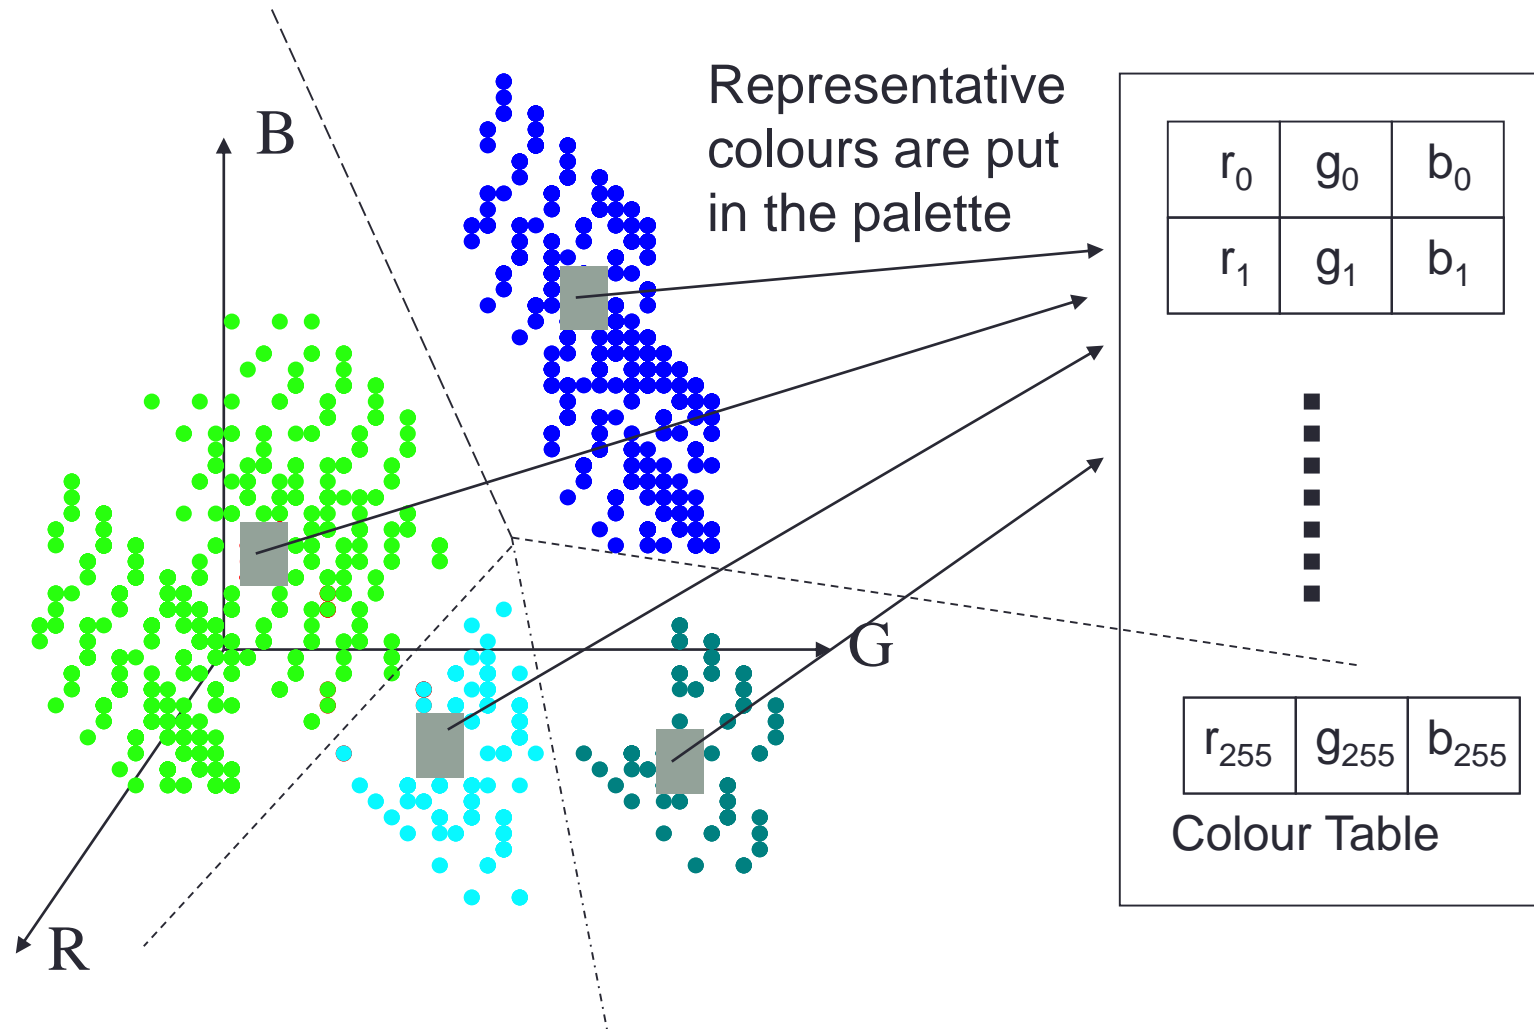
Building a Palette



Group pixels that are close to each other, and replace them by one single colour

Close to each other means of a similar colour

Building a Palette



Building a Palette

- Many clustering algorithms exist
 - supervised
 - unsupervised
- We know how many clusters we need: one per palette entry
- We need clusters that are spread across the colour space
- A supervised method....
- K-Means Clustering
 - Start with estimates of the mean of each cluster
$$\mu_1, \mu_2, \dots, \mu_k$$
 - Assign each point, p , to the cluster where
$$|p - \mu_i|$$
is smallest
 - Recompute the means
 - Repeat until no changes are made to the clusters



Evaluating Compression

- **Fidelity Criteria:** success is judged by comparing original and compressed versions
- Some measures are objective, e.g. root **mean square error** (e_{rms}) and **signal to noise ratio (SNR)**
- Let $f(x,y)$ be the input image, $f'(x,y)$ be reconstructed input image from compressed bit stream, then

$$e_{rms} = \left(\frac{1}{MN} \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} (f'(x,y) - f(x,y))^2 \right)^{1/2}$$

$$SNR = \frac{\sum_{x=0}^{M-1} \sum_{y=0}^{N-1} (f'(x,y))^2}{\sum_{x=0}^{M-1} \sum_{y=0}^{N-1} (f'(x,y) - f(x,y))^2}$$

Fidelity Criteria



$$\begin{array}{ll} e_{rms} = 6.93 & e_{rms} = 6.78 \\ SNR_{rm} = 10.25 & SNR_{rm} = 10.39 \end{array}$$

Fidelity Criteria

- e_{rms} and SNR are convenient objective measures
- Most decompressed images are viewed by human beings
- Subjective evaluation of compressed image quality by human observers are often more appropriate

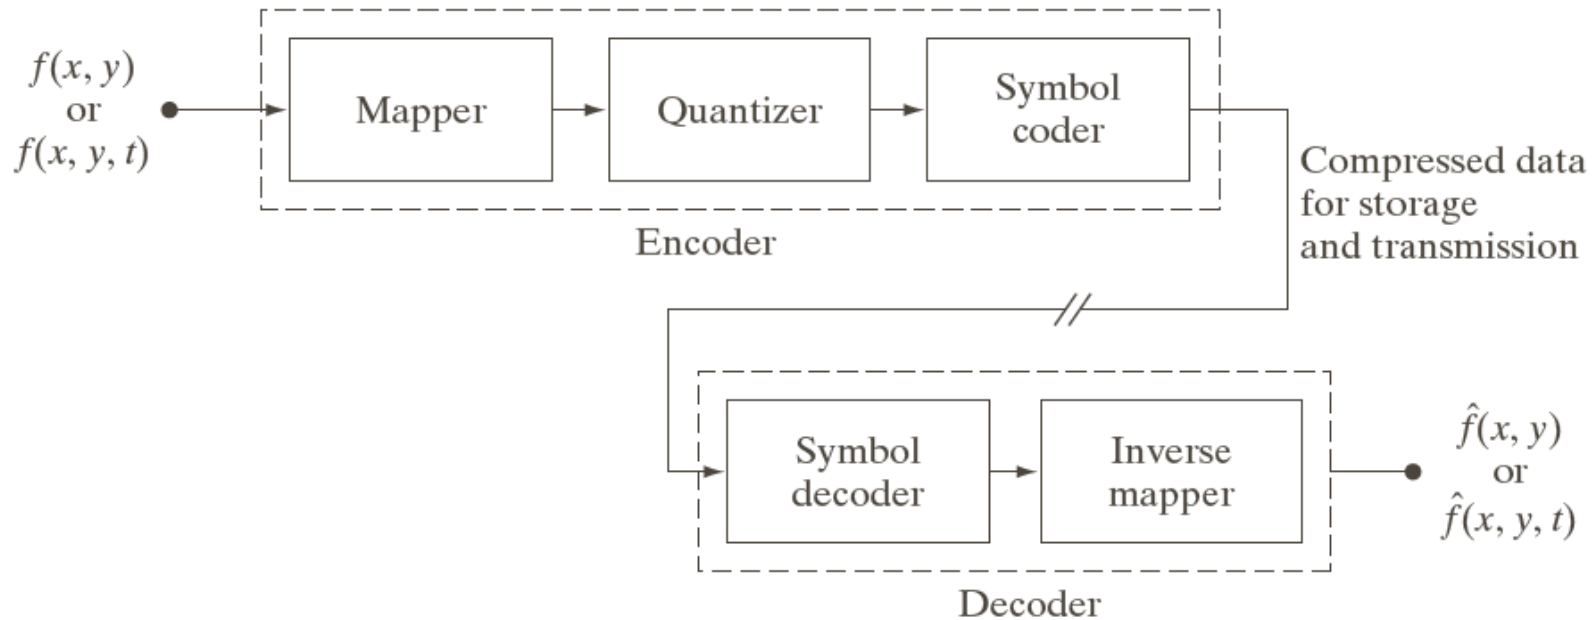
Value	Rating	Description
1	Excellent	An image of extremely high quality, as good as you could desire.
2	Fine	An image of high quality, providing enjoyable viewing. Interference is not objectionable.
3	Passable	An image of acceptable quality. Interference is not objectionable.
4	Marginal	An image of poor quality; you wish you could improve it. Interference is somewhat objectionable.
5	Inferior	A very poor image, but you could watch it. Objectionable interference is definitely present.
6	Unusable	An image so bad that you could not watch it.

Rating scale of the Television Allocations Study Organization. (Freundendall and Behrend.)

Image Compression Systems

- Mapper
 - transforms input data in a way that facilitates reduction of interpixel redundancies
 - reversible
- Quantiser
 - transforms input data in a way that facilitates reduction of psychovisual redundancies
 - not reversible
- Symbol coder
 - assigns the shortest code to the most frequently occurring output values
 - reversible

Image Compression Systems



Functional block diagram of a general image compression system

Image Compression in Matlab

- Imwrite provides control of compression methods and parameters
- `[peaksnr, snr] = psnr(image, reference image)` supports evaluation



JPEG quality 25/100
SNR = 30.0



JPEG quality 100/100
SNR = 53.8

Next Week

Conclusion and Exam Revision